

HarvardX Data Science Professional Certificate: Capstone Project

Delong Meng

12/4/2019

Introduction

Recommendation system represents a classical application of machine learning technology. For example, for a movie recommendation system, the goal is to predict how a given user will rate a specific movie based on how the user rate other movies and how the movie is rated by other users. In this project, I will combine several machine learning strategies to construct a movie recommendation system based on the “MovieLens” dataset. The full MovieLens dataset, which can be found here: <https://grouplens.org/datasets/movielens/latest/>, is rather large. To make the computation easier, in this project I use the “10M” version of MovieLens dataset instead (<https://grouplens.org/datasets/movielens/10m/>). This 10M MovieLens dataset has around 10 million ratings and 100,000 tag applications applied to 10,000 movies by 72,000 users.

Methods

I first downloaded the dataset and organize it into a format that is ready to explore. There are several features of this dataset, including movie ID, user ID, movie title (including the year the movie was released), genres, rating, and the timestamp of the rating. I extracted the release year from the title, and the rating year, and then calculated the age of the movie at the time of the rating. For the genres information, because it is a combination of different classifications, I split one single combination into individual records for further exploration.

To train the machine learning algorithm, I first randomly divided the 10M MovieLens dataset into a training set called “edx” (90%), and a test set called “validation” (10%). The “validation” set does not contain any user or movie that are absent in the “edx” set.

To evaluate the performance of the algorithm, I use Root Mean Square Error (RMSE) as an indicator. Only the “edx” set is used to develop the algorithm, and the “validation” set is used to get the RMSE value. I will first build a baseline prediction model including age effect, movie effect, user effect, and regularization of movie effect and user effect. I will evaluate those model and choose the best one to continue with using matrix factorization technique. More details can be found in the Results section.

1. Installing essential packages

I first set up the working environment by installing essential packages:

```
install.packages("tidyverse")
install.packages("caret")
install.packages("data.table")
install.packages("lubridate")
install.packages("recosystem")
install.packages("kableExtra")
```

```
library(tidyverse)
library(caret)
library(data.table)
library(lubridate)
library(stringr)
library(recosystem)
library(kableExtra)
library(tinytex)
```

2. Data downloading and preparation

Information of the MovieLens 10M dataset can be found: <https://grouplens.org/datasets/movielens/10m/>

The dataset can be downloaded here: <http://files.grouplens.org/datasets/movielens/ml-10m.zip>

I use the following code to download the dataset and combine the information of the ratings and the movies into *movielens*:

```
dl <- tempfile()
download.file("http://files.grouplens.org/datasets/movielens/ml-10m.zip", dl)
ratings <- fread(text = gsub(":", "\t", readLines(unzip(dl, "ml-10M100K/ratings.dat"))),
  col.names = c("userId", "movieId", "rating", "timestamp"))
movies <- str_split_fixed(readLines(unzip(dl, "ml-10M100K/movies.dat")), "\\:", 3)
colnames(movies) <- c("movieId", "title", "genres")
movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(levels(movieId))[movieId],
  title = as.character(title),
  genres = as.character(genres))
movielens <- left_join(ratings, movies, by = "movieId")
```

Now we split the MovieLens dataset into Training (*edx*) and Validation (*validation*) sets. The Validation set will be 10% of MovieLens data.

To successfully perform validation, we need to make sure all *userIds* and *movieIds* in the *validation* set are also in the *edx* set.

```
set.seed(1, sample.kind="Rounding")
test_index <- createDataPartition(y = movielens$rating, times = 1, p = 0.1, list = FALSE)
edx <- movielens[-test_index,]
temp <- movielens[test_index,]
validation <- temp %>%
  semi_join(edx, by = "movieId") %>%
  semi_join(edx, by = "userId")
removed <- anti_join(temp, validation)
edx <- rbind(edx, removed)
rm(dl, ratings, movies, test_index, temp, movielens, removed)
```


3. Data exploration

3.1 Overall profile of the dataset

Let's first have a general overview of the dataset:

```
head(edx)
```


```
##      userId movieId rating timestamp                title
## 1         1    122      5 838985046          Boomerang (1992)
## 2         1    185      5 838983525            Net, The (1995)
## 3         1    292      5 838983421            Outbreak (1995)
## 4         1    316      5 838983392            Stargate (1994)
## 5         1    329      5 838983392 Star Trek: Generations (1994)
## 6         1    355      5 838984474    Flintstones, The (1994)
##
##                genres
## 1          Comedy|Romance
## 2      Action|Crime|Thriller
## 3 Action|Drama|Sci-Fi|Thriller
## 4      Action|Adventure|Sci-Fi
## 5 Action|Adventure|Drama|Sci-Fi
## 6    Children|Comedy|Fantasy
```



We can see the overall distribution of all of the ratings. It is skewed to the right. All half stars are less frequent than full stars. A red dashed line of the overall average rating is also plotted here as a reference.

```
dim(edx) # 9000055      6
n_distinct(edx$movieId) # 10677
n_distinct(edx$title) # 10676: there might be movies of different IDs with the same title
n_distinct(edx$userId) # 69878
n_distinct(edx$movieId)*n_distinct(edx$userId) # 746087406
n_distinct(edx$movieId)*n_distinct(edx$userId)/dim(edx)[1] # 83
```

As shown above, this edx dataset has 10677 distinct movies and 69878 distinct users. If every user rated on every movie, we would have $10677 \times 69878 = 746087406$ ratings. However, we only have 9000055 ratings, which is only 1/83 of the number of all possible ratings. We can visualize 100 random samples of users and 100 random samples of movies to see how sparse this dataset is, which gives us an idea why this recommendation system is such a challenging task:



3.2 Extracting age of movies at rating

Every movie was released in a certain year, which is provided in the *title* of the movie. Every user rated a movie in a certain year, which is included in the *timestamp* information. I define the difference between these two years, i.e., how old the movie was when it was watched/rated by a user, as the **age of movies at rating**. From the original dataset, I first extracted the rating year (*year Rated*) from *timestamp*, and then extracted the release year (*year Released*) of the movie from the *title*. *age_at_rating* was later calculated.

```
# convert timestamp to year
edx_1 <- edx %>% mutate(year Rated = year(as_datetime(timestamp)))
# extract the release year of the movie
# edx_1 has year Rated, year Released, age_at_rating, and titles without year information
edx_1 <- edx_1 %>% mutate(title = str_replace(title, "^(.+)\s\\((\\d{4})\\)$", "\\1__\\2" )) %>%
  separate(title, c("title", "year Released"), "__") %>%
  select(-timestamp)
edx_1 <- edx_1 %>% mutate(age_at_rating = as.numeric(year Rated) - as.numeric(year Released))
head(edx_1)
```

```
##   userId movieId rating      title year_released
## 1      1     122      5    Boomerang      1992
## 2      1     185      5      Net, The      1995
## 3      1     292      5    Outbreak      1995
## 4      1     316      5    Stargate      1994
## 5      1     329      5 Star Trek: Generations 1994
## 6      1     355      5 Flintstones, The      1994
##                                     genres year Rated age_at_rating
## 1                                     Comedy|Romance      1996      4
## 2                                     Action|Crime|Thriller      1996      1
## 3 Action|Drama|Sci-Fi|Thriller      1996      1
## 4 Action|Adventure|Sci-Fi      1996      2
## 5 Action|Adventure|Drama|Sci-Fi      1996      2
## 6 Children|Comedy|Fantasy      1996      2
```

3.3 Extracting the genres information

The genres information was provided in the original dataset as a combination of different classifications. For example (see above output), the movie “Boomerang” (movieId 122) was assigned “Comedy|Romance”, and “Flintstones, The” (movieId 355) is “Children|Comedy|Fantasy”. Both are combinations of different ones, while they actually share one genre (Comedy). It’ll make more sense if we first split these combinations into single ones:

```
# edx_2: the mixture of genres is split into different rows
edx_2 <- edx_1 %>% separate_rows(genres, sep = "\\|") %>% mutate(value=1)
n_distinct(edx_2$genres) # 20: there are 20 different types of genres
genres_rating <- edx_2 %>% group_by(genres) %>% summarize(n=n())
genres_rating
```

```
## [1] 20
```

```
## # A tibble: 20 x 2
##   genres      n
##   <chr>    <int>
## 1 (no genres listed)      7
## 2 Action      2560545
## 3 Adventure  1908892
## 4 Animation   467168
## 5 Children    737994
## 6 Comedy     3540930
## 7 Crime      1327715
## 8 Documentary   93066
## 9 Drama      3910127
## 10 Fantasy     925637
## 11 Film-Noir   118541
## 12 Horror      691485
## 13 IMAX        8181
## 14 Musical     433080
## 15 Mystery     568332
## 16 Romance    1712100
## 17 Sci-Fi     1341183
## 18 Thriller   2325899
## 19 War        511147
```

```
## 20 Western          189394
```

We get the information of 20 different types of genres, and numbers of movie ratings in each type. Note that the first type is “(no genres listed)”, which is not really a genre, but just reflects the fact that for 7 ratings (of one single movie), genres info was not provided.

We can figure out the only one movie with no genres information is movieId 8606 titled “Pull My Daisy”, released in 1958:

```
## # A tibble: 1 x 5
## # Groups:   movieId, title, year_released [1]
##   movieId title          year_released genres          n
##   <int> <chr>          <chr>          <fct>          <int>
## 1     8606 Pull My Daisy 1958          (no genres listed)      7
```

Splitting the genres information into multiple row can facilitate the exploration of genres. However, one thing to keep in mind is, if we consider one row as one record (here for our movie recommendation system every rating for a certain movie rated by a certain user should be one record), the above transformation of the dataset actually duplicated each record into multiple ones, depending on the combination of the genres for each movie.

To avoid this problem and make more sense if we want to utilize the genres information in building the prediction model, genres of each movie (also each record) should be split into multiple columns to indicate different combinations of the 19 basic genres. We can achieve this goal by spreading genres to the “wide” format:

```
# edx_3 is the final version for exploration of the effects of movie year, age, rating year, and genres
edx_3 <- edx_2 %>% spread(genres, value, fill=0) %>% select(-(“no genres listed”))
dim(edx_3) # 9000055      26
```

```
## [1] 9000055      26
```

```
head(edx_3)
```

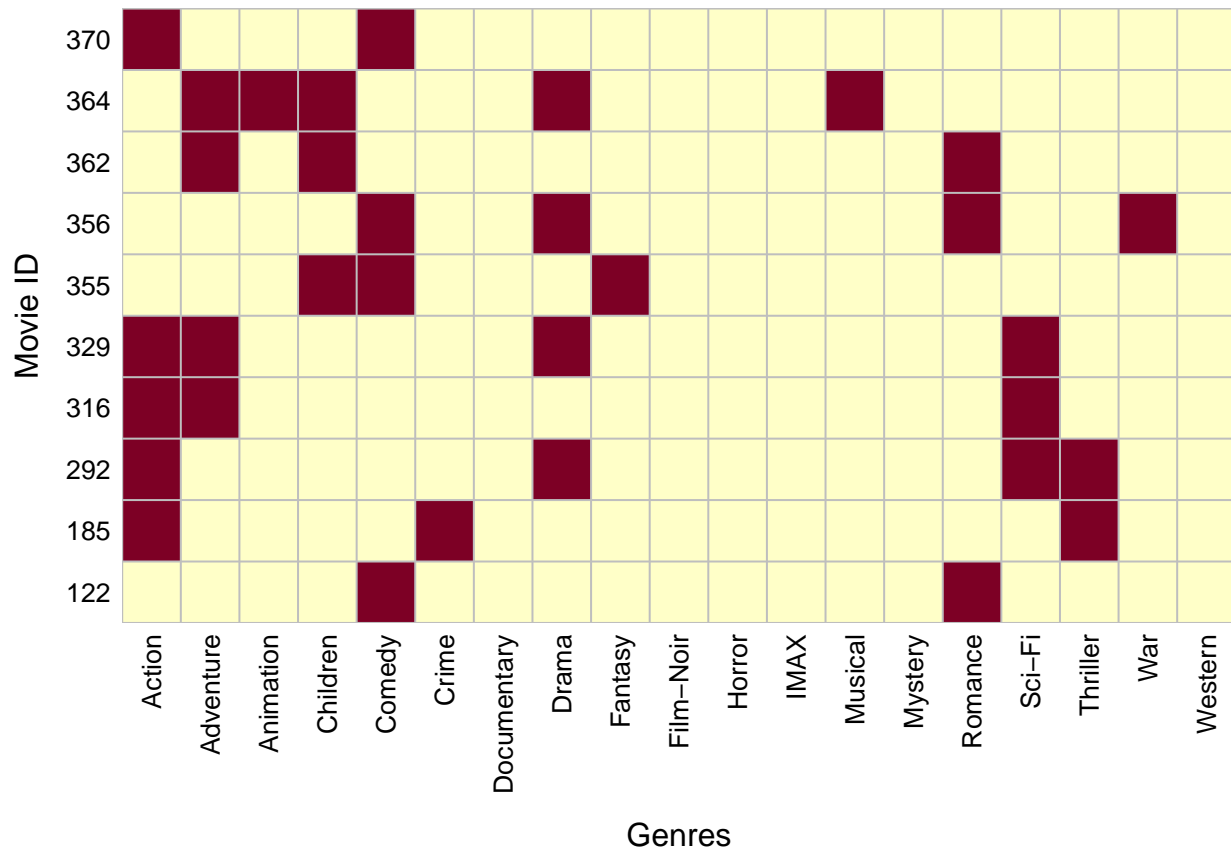
```
##   userId movieId rating          title year_released year_rated
## 1      1      122      5      Boomerang          1992          1996
## 2      1      185      5      Net, The          1995          1996
## 3      1      292      5      Outbreak          1995          1996
## 4      1      316      5      Stargate          1994          1996
## 5      1      329      5 Star Trek: Generations 1994          1996
## 6      1      355      5      Flintstones, The 1994          1996
##   age_at_rating Action Adventure Animation Children Comedy Crime
## 1              4      0          0          0          0      1      0
## 2              1      1          0          0          0      0      1
## 3              1      1          0          0          0      0      0
## 4              2      1          1          0          0      0      0
## 5              2      1          1          0          0      0      0
## 6              2      0          0          0          1      1      0
##   Documentary Drama Fantasy Film-Noir Horror IMAX Musical Mystery Romance
## 1              0      0          0          0      0      0      0      0      1
## 2              0      0          0          0      0      0      0      0      0
## 3              0      1          0          0      0      0      0      0      0
## 4              0      0          0          0      0      0      0      0      0
```

```

## 5      0      1      0      0      0      0      0      0      0
## 6      0      0      1      0      0      0      0      0      0
##  Sci-Fi Thriller War Western
## 1      0      0      0      0
## 2      0      1      0      0
## 3      1      1      0      0
## 4      1      0      0      0
## 5      1      0      0      0
## 6      0      0      0      0

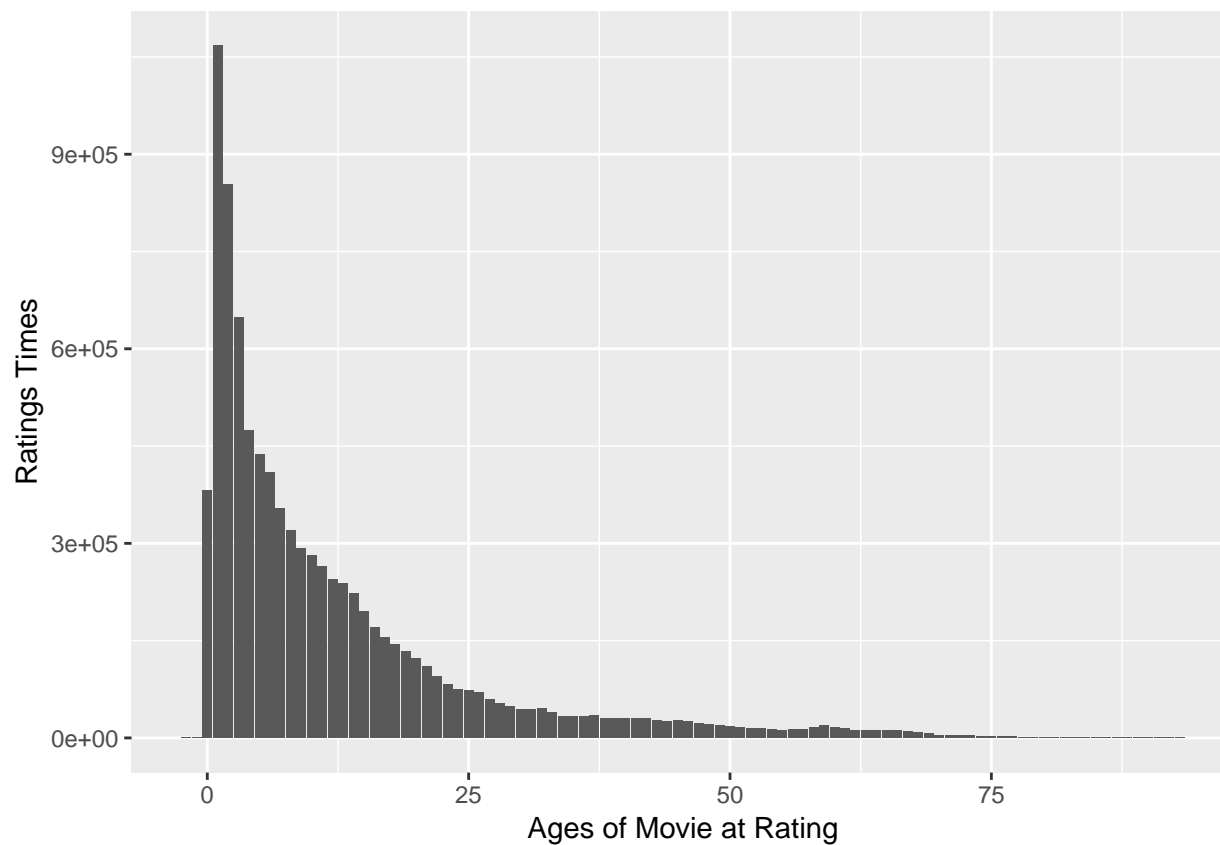
```

By doing this, we can actually visualize the genres of each movie. For example, each row represents a movie, and each column represents a genre. Each movie can have a unique combination of different genres. Here we look at the first 10 movies:




3.4 Exploration of age effect

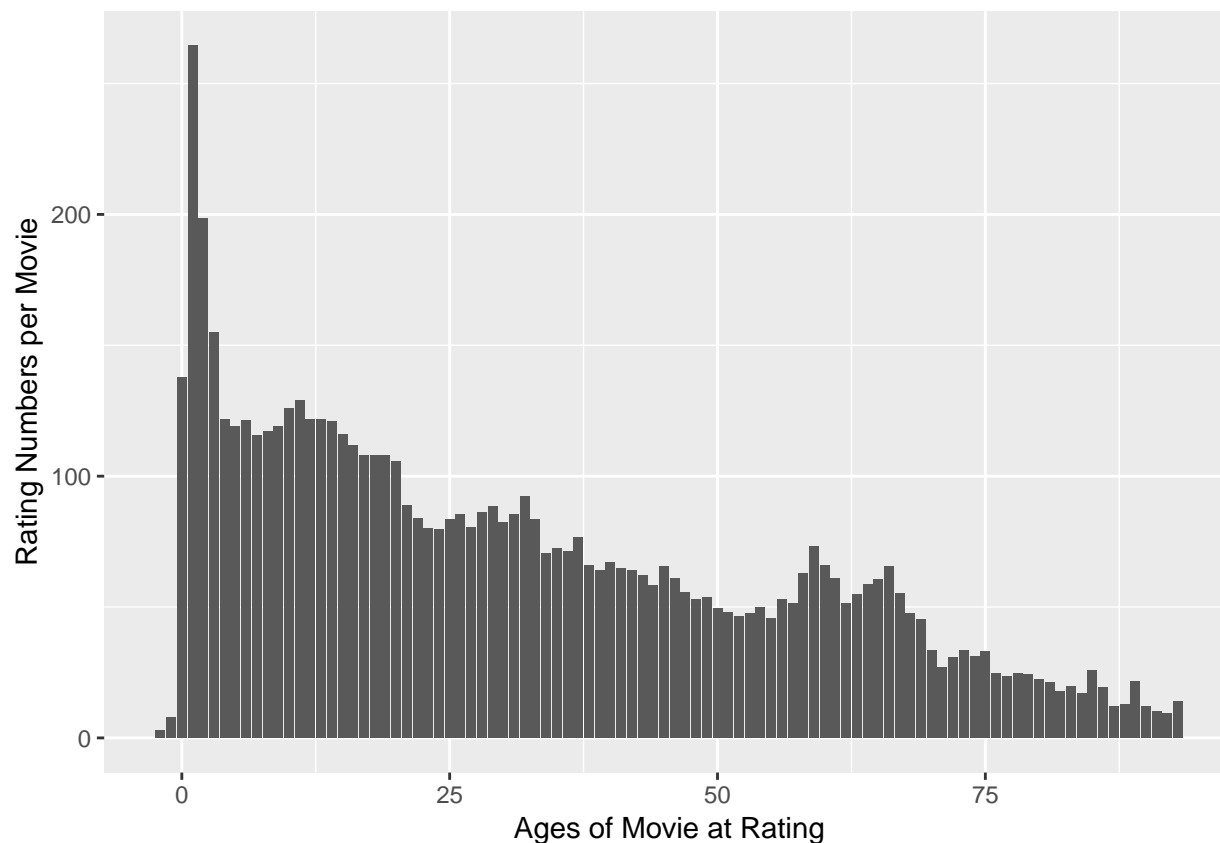
Could how old/new a movie was when watched/rated by a user affect the rating it got? First, let's take a look at the distribution of rating numbers according to the age of movies at the time of rating.



It appears that mosting ratings were given for newer movies (within 20 years old at the time of rating). However, is it simply because there are more newer movies in the whole dataset that have been rated by users? Let's take a look at the distribution of movie numbers:




We see that actually as you can imagine, we have a similar trend that most movies in the dataset are within around 25 years old when being rated. To have a better idea of if newer movies get rated more frequently than older movies, we'd better look at rating times per movie for a given age group:

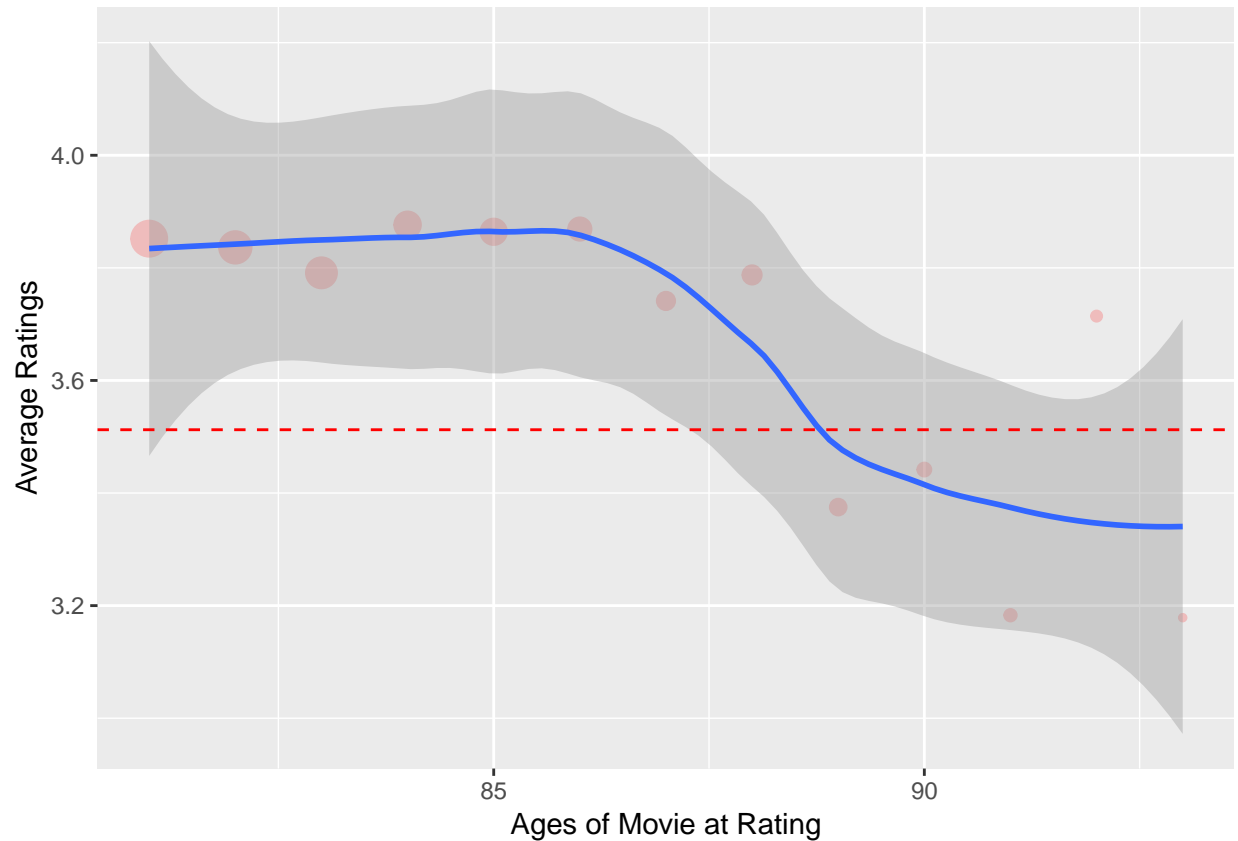


Although the trend is still that older movies get less frequently rated. It's not as dramatic as the overall rating times when we normalize by the movie numbers now.

The next question we try to address here is whether the age of movie at the time of rating affects rating? We can plot the average ratings for movies with a certain age at the time of ratings against the age. The red circles represent the average ratings while the size of a circle corresponds to the numbers of ratings for the movies with the same age at the time of rating. The blue line shows a smoothened overall trend. I also add a red dashed line here to represent the overall average rating. I further zoom in the critical area when the ratings get close to the average.








Interestingly, movies with less than 18 years tend to have around average or slightly lower than average ratings, while older movies tend to have much greater than average ratings (with only handful exceptions).

This could be because old movies that are still widely watched are usually those with good reputation and are widely recommended. In another word, they have been already selected. On the other hand, newer movies have not been judged enough by customers to have a selection effect.

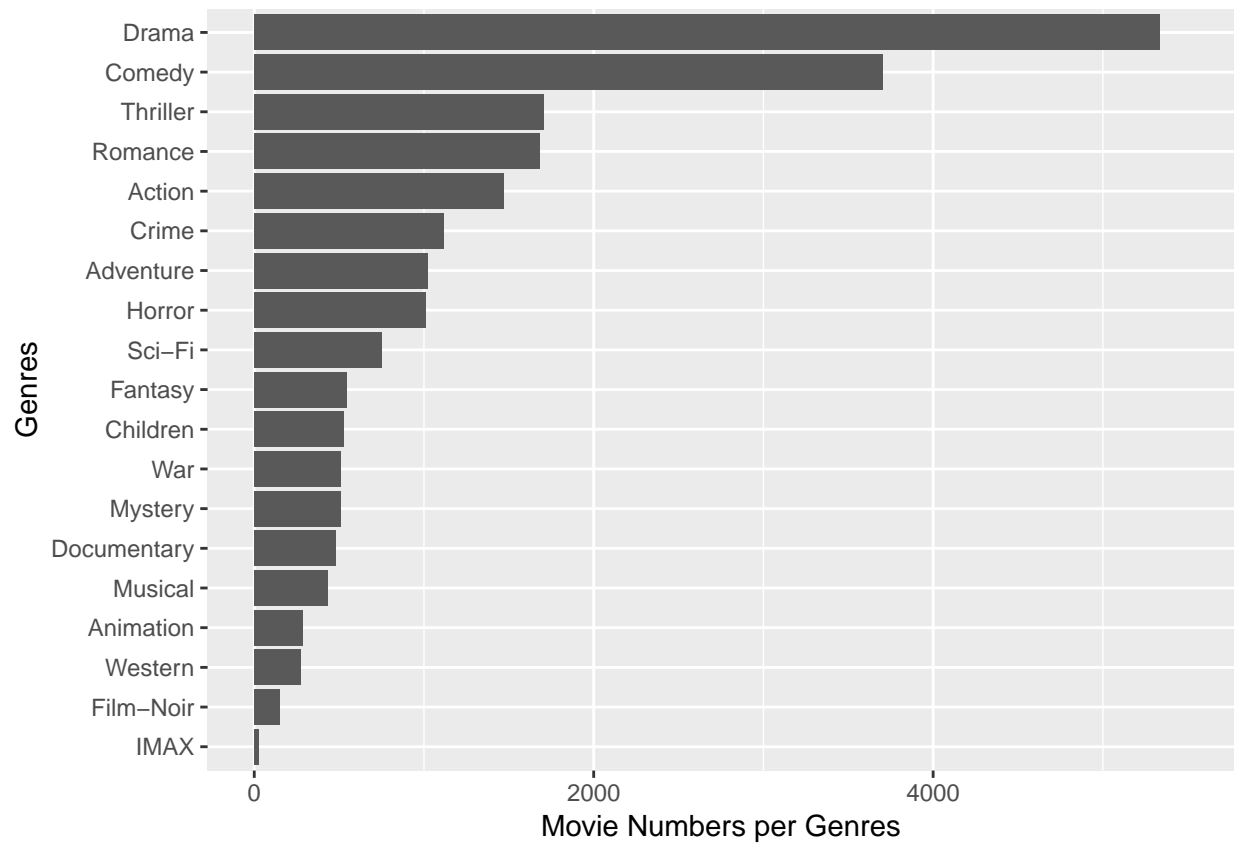
3.5 Exploration of genres

Next let's explore the effect of genres on movie ratings. First, are movies of different genres rated in different frequencies?



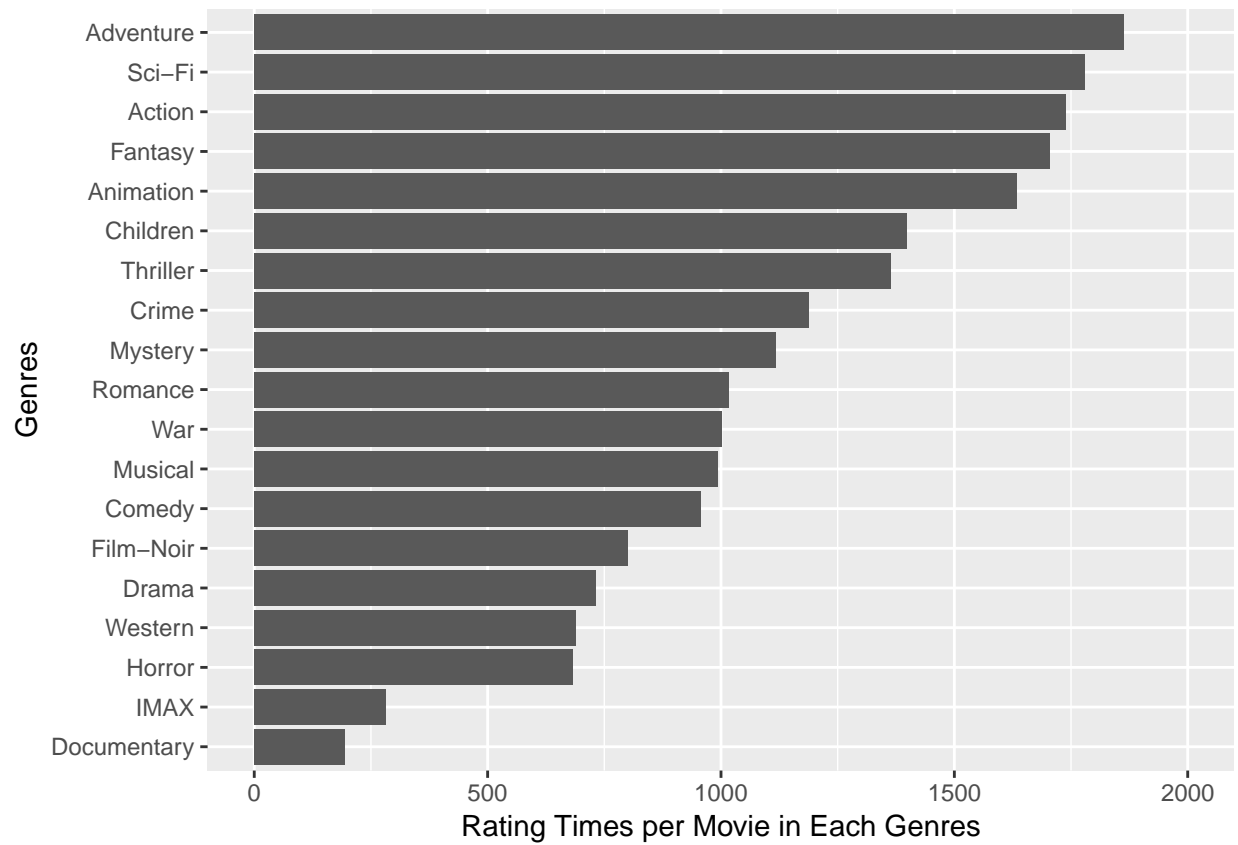
This figure shows the number of movies in each genres. Note that one single movie can actually belong to different genres. The Drama genres is rated the most while the IMAX genres is the least rated. However, this doesn't necessarily mean people prefer to rate the Drama movies over other types, because this could simply reflect there are more movies in the Drama genres.

Let's check the numbers of movies of each genres:



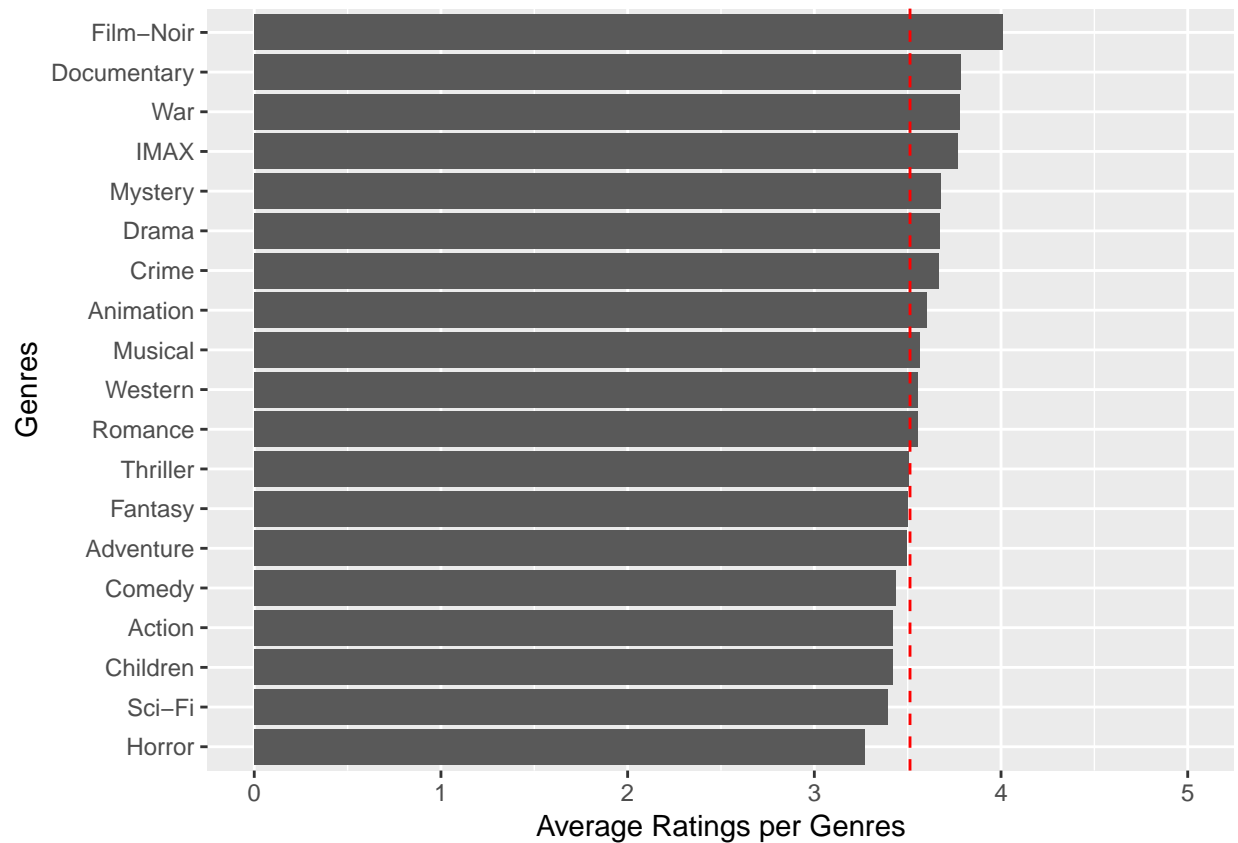
As expected, we see the same trend as we saw in total rating numbers.

So, to evaluate really how genres affect rating frequencies of movies, let's check the average numbers of ratings per movie in each genres:



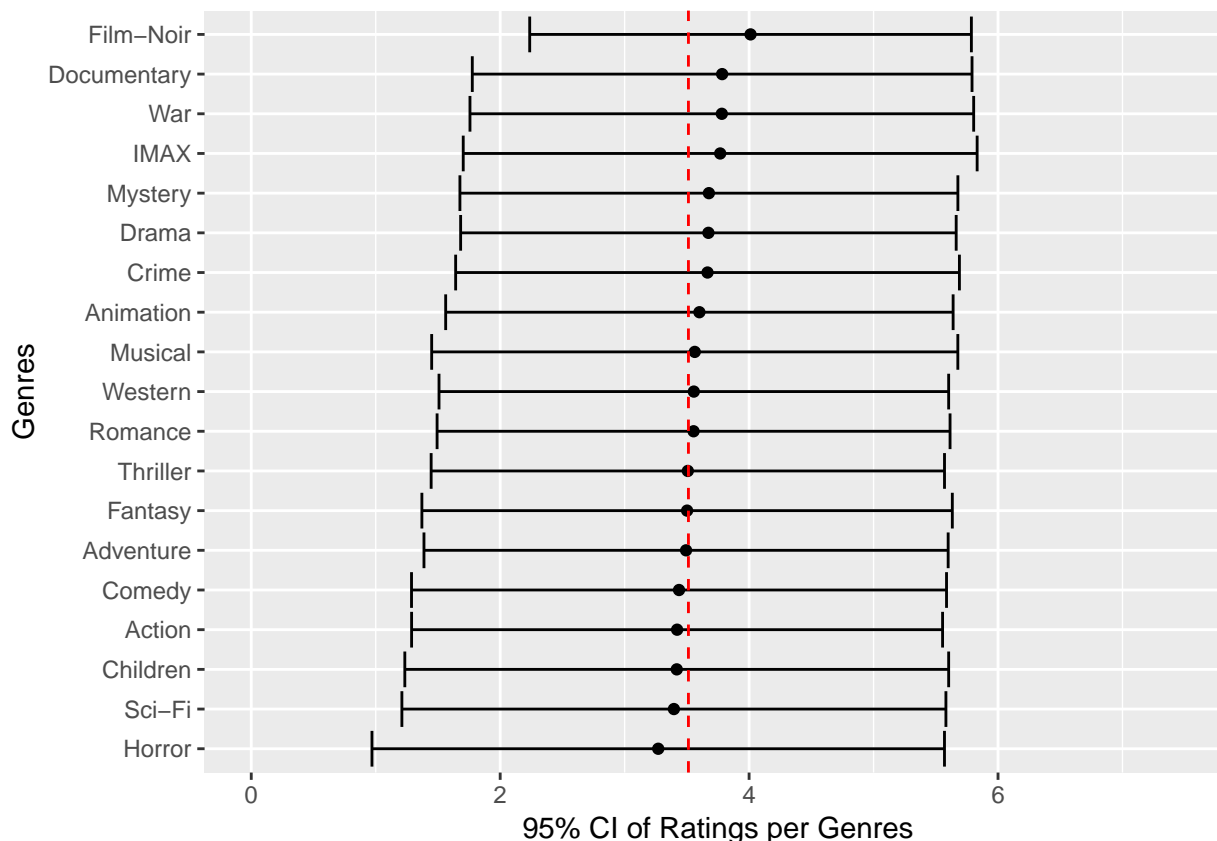
Now we see a different trend. Interestingly, Adventure and Sci-Fi are more tented to be rated although they don't include as many movies compared to Drama.

Next question is: what about the value of ratings themselves (which we actually care about)?



Now we find that some genres tend to have higher ratings than the average (such as Film-Noir) and some tend to have lower ratings (such as Horror). However, overall, the genres effect seems to be rather minor.

Another way to visualize the data is: instead of only looking at average, we can also see the distribution. So we calculate both mean and sd of each genres and generate the 95% CI:



We can also see that the genres only slightly affect movie ratings.

4. Modeling strategy

Based on the exploration of the data described above, age of movie at rating seems to affect the rating, while genres does not add much information. Also, the effect of genres could also be included in the movie effect itself. Therefore, in the next section, to first build a baseline prediction model, I will consider the effects of movie age at rating, movie effect, and user effect. Regularization of movie effect and user effect will also be used to build a more robust model. I will evaluate these models and choose the best to go with. Residuals will be calculated and used as the input of matrix factorization technique.

The baseline model was generated following the instructions described in the “Recommendation Systems” of the text book (<https://rafalab.github.io/dsbook/large-datasets.html#recommendation-systems>). For example, the movie and user effect model describes the rating $Y_{u,i} = \mu + b_i + b_u + \text{error}$, while μ is the average of all ratings, b_i and b_u are item (i.e., movie) bias and user bias, respectively. To avoid over-training caused by estimates from small sample sizes, I use regularization to add penalties to shrink the estimates from small samples sizes.

As mentioned in the text book (<https://rafalab.github.io/dsbook/large-datasets.html#recommendation-systems>), matrix factorization is a very useful technique used in recommendation systems. It can identify the similar patterns of movies as well as users in terms of being rated or rating a movie. I first calculated the residual ($r_{u,i} = y_{u,i} - \mu - b_i - b_u$) of the predictions based on the baseline model, and used the matrix factorization strategy to model the residual. Matrix factorization tries to decompose the rating matrix into user matrix p_u and item (movie) matrix q_i , so that $r_{u,i}$ can be explained by $p_u * q_i$. To achieve this goal, I used the R package “recosystem” (<https://cran.r-project.org/web/packages/recosystem/vignettes/introduction.html>), which can conveniently apply matrix factorization to my dataset using the parallel stochastic gradient descent algorithm.

This “recoSystem” R package is a wrapper of an open source C++ library LIMBF. A good thing about “recoSystem” is that it can significantly reduce memory use by storing input, model and output information in the hard disk instead of using memory.

Results

Define RMSE: residual mean squared error

```
RMSE <- function(true_ratings, predicted_ratings){  
  sqrt(mean((true_ratings - predicted_ratings)^2))  
}
```

Model 1

First model: use average ratings for all movies regardless of user

In the first model, just based on the ratings itself, to minimize the RMSE, the best prediction of ratings for each movie will be the overall average of all ratings. The average rating is $\mu = 3.51247$, and the naive RMSE is 1.0612.

```
mu <- mean(edx$rating)  
mu
```

```
## [1] 3.51247
```

```
naive_rmse <- RMSE(validation$rating, mu)  
naive_rmse
```

```
## [1] 1.0612
```

```
rmse_results <- data_frame(Model = "Just the average", RMSE = naive_rmse)  
rmse_results
```


```
## # A tibble: 1 x 2  
##   Model      RMSE  
##   <chr>      <dbl>  
## 1 Just the average 1.06120
```

Model 2

Modeling Age Effects: adding b_a to represent ratings on movies with certain age

Because earlier we saw that the age of movies at the time of rating seems to affect ratings, I try to see if add a bias of age (b_a) to the model could better predict the ratings. First let's calculate the age bias and take a look at its distribution. Then we will make predictions and evaluate the RMSE using the *validation* set.

```
age_effect<- edx_1 %>%
  group_by(age_at_rating) %>%
  summarize(b_a = mean(rating)-mu)
age_effect %>% qplot(b_a, geom="histogram", bins = 10, data = ., color = I("black"))
```



```
validation_1 <- validation %>%
  mutate(year_rated = year(as_datetime(timestamp)))%>%
  mutate(title = str_replace(title,"^(.+)"\\s\\((\\d{4})\\)"$, "\\1__\\2" )) %>%
  separate(title,c("title","year_released"),"__") %>%
  select(-timestamp) %>%
  mutate(age_at_rating= as.numeric(year_rated)-as.numeric(year_released))

predicted_ratings_2 <- mu + validation_1 %>%
  left_join(age_effect, by='age_at_rating') %>%
  pull(b_a)
model_2_rmse <- RMSE(validation$rating,predicted_ratings_2) # 1.05239
rmse_results <- bind_rows(rmse_results,
  data_frame(Model="Age Effect Model",
    RMSE = model_2_rmse))

rmse_results
```

```
## # A tibble: 2 x 2
##   Model      RMSE
##   <chr>      <dbl>
## 1 Just the average 1.06120
```

```
## 2 Age Effect Model 1.05239
```


We can see that Age Effect Model did not improve the RMSE much. For this reason we will stop using the ages of movies as a predictor.

Model 3

Modeling movie effects: adding b_i to represent average ranking for movie_i

Since the intrinsic features of a movie could obviously affect the ratings of a movie, we add the bias of movie/item (b_i) to the model, i.e., for each movie, the average of the ratings on that specific movie will have a difference from the overall average rating of all movies. We can plot the distribution of the bias and calculate the RMSE of this model.

```
movie_avgs <- edx %>%  
  group_by(movieId) %>%  
  summarize(b_i = mean(rating - mu))  
movie_avgs %>% qplot(b_i, geom = "histogram", bins = 10, data = ., color = I("black"))
```



```
predicted_ratings_3 <- mu + validation %>%  
  left_join(movie_avgs, by = 'movieId') %>%  
  pull(b_i)  
model_3_rmse <- RMSE(validation$rating, predicted_ratings_3)  
rmse_results <- bind_rows(rmse_results,
```

```

data_frame(Model="Movie Effect Model",
            RMSE = model_3_rmse))
rmse_results

```

```

## # A tibble: 3 x 2
##   Model          RMSE
##   <chr>          <dbl>
## 1 Just the average 1.06120
## 2 Age Effect Model 1.05239
## 3 Movie Effect Model 0.943909

```

Adding the movie bias successfully brought the RMSE to lower than 1.

Model 4

User effects: adding `b_u` to represent average ranking for `user_u`

Similar to the movie effect, intrinsic features of a given user could also affect the ratings of a movie. For example, a stricter user could give lower scores for all movies he/she watched than rated by other users. We now further add the bias of user (`b_u`) to the movie effect model.

```

user_avgs <- edx %>%
  left_join(movie_avgs, by='movieId') %>%
  group_by(userId) %>%
  summarize(b_u = mean(rating - mu - b_i))
predicted_ratings_4 <- validation %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  mutate(pred = mu + b_i + b_u) %>%
  pull(pred)
model_4_rmse <- RMSE(validation$rating,predicted_ratings_4)
rmse_results <- bind_rows(rmse_results,
                          data_frame(Model="Movie + User Effects Model",
                                      RMSE = model_4_rmse))
rmse_results

```

```

## # A tibble: 4 x 2
##   Model          RMSE
##   <chr>          <dbl>
## 1 Just the average 1.06120
## 2 Age Effect Model 1.05239
## 3 Movie Effect Model 0.943909
## 4 Movie + User Effects Model 0.865349

```

Adding the user effect dramatically increased the RMSE to lower than 0.9.

Model 5

Regularization of movie effect: control the total variability of the movie effects

A machine learning model could be over-trained if some estimates were from a very small sample size. Regularization technique should be used to take into account the number of ratings made for a specific

movie, by adding a larger penalty to estimates from smaller samples. To do this, a parameter λ will be used. Cross validation within the test set can be performed to optimize this parameter before being applied to the validation set.

1. Perform cross validation to determine the parameter λ

To train the parameter λ , I use 10-fold cross validation here within only the *edx* set, because the *validation* set should not be used to train any parameter.

Specifically, I first randomly split the training set (*edx*) into 10 parts. Each time, I combine 9 parts as a *train_set*, and use the 10th part as a *test_set*. I will make sure all *userId*s and *movieId*s in test set are also in the train set. For a given range of different values of λ , I will build the model using the *train_set* and evaluate the performance using the *test_set*. By doing this, I will get a set of RMSEs corresponding to these different λ s.

In total, there will be 10 possible combinations of the 10 subsets of *edx*, so I will have 10 sets of *train_set* and *test_set*. Thus, I will do 10 times of training and get 10 sets of RMSEs. For each λ value I try, I will have 10 RMSEs and take the average. Then I will determine the minimal RMSE and use the corresponding λ as the optimized λ for model building and performance evaluation in the *validation* set.

```
# use 10-fold cross validation to pick a lambda for movie effects regularization
# split the data into 10 parts
set.seed(2019, sample.kind = "Rounding")
cv_splits <- createFolds(edx$rating, k=10, returnTrain = TRUE)

# define a matrix to store the results of cross validation
rmsees <- matrix(nrow=10, ncol=51)
lambdas <- seq(0, 5, 0.1)

# perform 10-fold cross validation to determine the optimal lambda
for(k in 1:10) {
  train_set <- edx[cv_splits[[k]],]
  test_set <- edx[-cv_splits[[k]],]

  # Make sure userId and movieId in test set are also in the train set
  test_final <- test_set %>%
    semi_join(train_set, by = "movieId") %>%
    semi_join(train_set, by = "userId")

  # Add rows removed from validation set back into edx set
  removed <- anti_join(test_set, test_final)
  train_final <- rbind(train_set, removed)

  mu <- mean(train_final$rating)
  just_the_sum <- train_final %>%
    group_by(movieId) %>%
    summarize(s = sum(rating - mu), n_i = n())


  rmsees[k,] <- sapply(lambdas, function(l){
    predicted_ratings <- test_final %>%
      left_join(just_the_sum, by='movieId') %>%
      mutate(b_i = s/(n_i+1)) %>%
      mutate(pred = mu + b_i) %>%
      pull(pred)
```

```

    return(RMSE(predicted_ratings, test_final$rating))
  })
}

rmse_cv <- colMeans(rmses)
qplot(lambdas, rmse_cv)
lambda <- lambdas[which.min(rmse_cv)] #2.2

```



From the 10-fold cross validation, we get an optimized value of lambda: 2.2.

2. Model generation and prediction

Regularized Movie Effect Model

Using the optimized lambda, we can now perform prediction and evaluate the RMSE in the *validation* set.

```

mu <- mean(edx$rating)
movie_reg_avgs <- edx %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+lambda), n_i = n())
predicted_ratings_5 <- validation %>%
  left_join(movie_reg_avgs, by = "movieId") %>%
  mutate(pred = mu + b_i) %>%
  pull(pred)
model_5_rmse <- RMSE(predicted_ratings_5, validation$rating) # 0.943852 not too much improved

```



```
rmse_results <- bind_rows(rmse_results,
                          data_frame(Model="Regularized Movie Effect Model",
                                     RMSE = model_5_rmse))
rmse_results
```

```
## # A tibble: 5 x 2
##   Model                                RMSE
##   <chr>                                <dbl>
## 1 Just the average                    1.06120
## 2 Age Effect Model                   1.05239
## 3 Movie Effect Model                 0.943909
## 4 Movie + User Effects Model         0.865349
## 5 Regularized Movie Effect Model    0.943852
```

The application of regularization does not improve the RMSE too much.

Model 6

Regularization of both movie and user effects (use the same lambda for both movie and user effects)

1. Perform cross validation to determine the parameter lambda

Similar to the movie effect, now we perform regularization on both movie and user effects. Still using 10-fold cross validation as described above, we will train one single lambda value for both movie and user effects.

```
# define a matrix to store the results of cross validation
lambdas <- seq(0, 8, 0.1)
rmse_2 <- matrix(nrow=10, ncol=length(lambdas))
# perform 10-fold cross validation to determine the optimal lambda
for(k in 1:10) {
  train_set <- edx[cv_splits[[k]],]
  test_set <- edx[-cv_splits[[k]],]

  # Make sure userId and movieId in test set are also in the train set
  test_final <- test_set %>%
    semi_join(train_set, by = "movieId") %>%
    semi_join(train_set, by = "userId")

  # Add rows removed from validation set back into edx set
  removed <- anti_join(test_set, test_final)
  train_final <- rbind(train_set, removed)

  mu <- mean(train_final$rating)

  rmse_2[k,] <- sapply(lambdas, function(l){
    b_i <- train_final %>%
      group_by(movieId) %>%
      summarize(b_i = sum(rating - mu)/(n()+1))
    b_u <- train_final %>%
      left_join(b_i, by="movieId") %>%
      group_by(userId) %>%
```

```

    summarize(b_u = sum(rating - b_i - mu)/(n()+1))
  predicted_ratings <-
    test_final %>%
    left_join(b_i, by = "movieId") %>%
    left_join(b_u, by = "userId") %>%
    mutate(pred = mu + b_i + b_u) %>%
    pull(pred)
  return(RMSE(predicted_ratings, test_final$rating))
})
}

rmses_2
rmses_2_cv <- colMeans(rmses_2)
rmses_2_cv
qplot(lambdas, rmses_2_cv)
lambda <- lambdas[which.min(rmses_2_cv)] #4.9

```

From the 10-fold cross validation, we get an optimized value of lambda: 4.9.

2. Model generation and prediction

Regularized Movie Effect and User Effect Model

Now we use this parameter lambda to predict the validation dataset and evaluate the RMSE.

```

mu <- mean(edx$rating)
b_i_reg <- edx %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+lambda))
b_u_reg <- edx %>%
  left_join(b_i_reg, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n()+lambda))
predicted_ratings_6 <-
  validation %>%
  left_join(b_i_reg, by = "movieId") %>%
  left_join(b_u_reg, by = "userId") %>%
  mutate(pred = mu + b_i + b_u) %>%
  pull(pred)
model_6_rmse <- RMSE(predicted_ratings_6, validation$rating) # 0.864818
rmse_results <- bind_rows(rmse_results,
  data_frame(Model="Regularized Movie + User Effect Model",
    RMSE = model_6_rmse))
rmse_results

```

```

## # A tibble: 6 x 2
##   Model                                RMSE
##   <chr>                                <dbl>
## 1 Just the average                    1.06120
## 2 Age Effect Model                   1.05239
## 3 Movie Effect Model                 0.943909
## 4 Movie + User Effects Model         0.865349

```

```
## 5 Regularized Movie Effect Model          0.943852
## 6 Regularized Movie + User Effect Model 0.864818
```

Regularization slightly improved the prediction performance of the model.

Model 7

Regularization of movie and user effects: use different lambdas

Optimizing λ_u (user effect) with fixed λ_i (movie effect)

1. Perform cross validation to determine the parameter λ_u for a given λ_i

Instead of optimizing the same lambda for both user and movie effect, here I tried to fix the lambda for movie using the value we got in model 5 ($\lambda_i=2.2$), and optimize the lambda for user (λ_u).

```
# define a matrix to store the results of cross validation
lambda_i <- 2.2
lambdas_u <- seq(0, 8, 0.1)
rmse3 <- matrix(nrow=10, ncol=length(lambdas_u))

# perform 10-fold cross validation to determine the optimal lambda
for(k in 1:10) {
  train_set <- edx[cv_splits[[k]],]
  test_set <- edx[-cv_splits[[k]],]

  # Make sure userId and movieId in test set are also in the train set
  test_final <- test_set %>%
    semi_join(train_set, by = "movieId") %>%
    semi_join(train_set, by = "userId")

  # Add rows removed from validation set back into edx set
  removed <- anti_join(test_set, test_final)
  train_final <- rbind(train_set, removed)

  mu <- mean(train_final$rating)

  rmse3[k,] <- sapply(lambdas_u, function(l){
    b_i <- train_final %>%
      group_by(movieId) %>%
      summarize(b_i = sum(rating - mu)/(n()+lambda_i))
    b_u <- train_final %>%
      left_join(b_i, by="movieId") %>%
      group_by(userId) %>%
      summarize(b_u = sum(rating - b_i - mu)/(n()+1))
    predicted_ratings <-
      test_final %>%
      left_join(b_i, by = "movieId") %>%
      left_join(b_u, by = "userId") %>%
      mutate(pred = mu + b_i + b_u) %>%
      pull(pred)
    return(RMSE(predicted_ratings, test_final$rating))
  })
}
```

```

}
rmse_3
rmse_3_cv <- colMeans(rmse_3)
rmse_3_cv
qplot(lambdas_u,rmse_3_cv)
lambda_u <- lambdas_u[which.min(rmse_3_cv)] #5

```

For a given λ_i of 2.2, we get an optimized λ_u of 5.

2. Model generation and prediction

Regularized Movie and User Effect Model with fixed λ for Movie Effect

Using the λ_i and λ_u we determined, I generated the prediction model and evaluated the RMSE in the validation set.

```

lambda_i <- 2.2
lambda_u <- 5
mu <- mean(edx$rating)
b_i_reg <- edx %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+lambda_i))
b_u_reg <- edx %>%
  left_join(b_i_reg, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n()+lambda_u))
predicted_ratings_7 <-
  validation %>%
  left_join(b_i_reg, by = "movieId") %>%
  left_join(b_u_reg, by = "userId") %>%
  mutate(pred = mu + b_i + b_u) %>%
  pull(pred)
model_7_rmse <- RMSE(predicted_ratings_7, validation$rating) # 0.86485
rmse_results <- bind_rows(rmse_results,
  data_frame(Model="Regularized Movie + User Effect Model Version 2",
    RMSE = model_7_rmse))
rmse_results

```

```

## # A tibble: 7 x 2
##   Model                                RMSE
##   <chr>                                <dbl>
## 1 Just the average                    1.06120
## 2 Age Effect Model                   1.05239
## 3 Movie Effect Model                 0.943909
## 4 Movie + User Effects Model         0.865349
## 5 Regularized Movie Effect Model     0.943852
## 6 Regularized Movie + User Effect Model 0.864818
## 7 Regularized Movie + User Effect Model Version 2 0.864850

```

While regularization using different parameters for user and item slightly improve RMSE (comparing Model 7 and Model 4), it did not improve the last regularization model (Model 6).

Model 8

Regularization of movie and user effects: use different lambdas

Optimizing λ_i (movie effect) with fixed λ_u (user effect)

1. Perform cross validation to determine the parameter λ_i for a given λ_u

Here I want to see if anything changes when I slightly change the strategy to fix λ_u and then choose λ_i . For the fixed $\lambda_u = 5$ (based on Model 7), I optimized λ_i using 10-fold cross validation and got a $\lambda_i = 4.6$.

2. Model generation and prediction

Regularized Movie and User Effect Model with fixed lambda for User Effect

A new model was generated similarly and RMSE determined using the *validation* set.

```
## # A tibble: 8 x 2
##   Model                                RMSE
##   <chr>                                <dbl>
## 1 Just the average                    1.06120
## 2 Age Effect Model                   1.05239
## 3 Movie Effect Model                 0.943909
## 4 Movie + User Effects Model         0.865349
## 5 Regularized Movie Effect Model     0.943852
## 6 Regularized Movie + User Effect Model 0.864818
## 7 Regularized Movie + User Effect Model Version 2 0.864850
## 8 Regularized Movie + User Effect Model Version 3 0.864819
```

The RMSE is slightly better than Model 7 and now comparable to Model 6, in which the same lambda was used for both user and item effects.

Model 9

Matrix Factorization based on the residuals of the baseline model

1. Best baseline model

Models 1-8 are all baseline models mainly based on movie effect and user effect. I compared the RMSEs and determined to go with “Regularized Movie + User Effect Model” (model 6) prior to further exploration because it gives the least RMSE.

2. Calculating the residuals

Next, we need to calculate the residuals. Although earlier we evaluated the models using the *validation* set, to calculate the residual we can’t use the *validation* set here, because the *validation* set can only be used at the end when evaluating the performance of the final model. Instead, we need to still use the training set *edx*. So how is the RMSE based on the *edx* set (training set)? As expected, it is slightly better than calculated using the *validation* set:

```

lambda <- 4.9
mu <- mean(edx$rating)
b_i_reg <- edx %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+lambda))
b_u_reg <- edx %>%
  left_join(b_i_reg, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n()+lambda))
predicted_ratings_6_edx <-
  edx %>%
  left_join(b_i_reg, by = "movieId") %>%
  left_join(b_u_reg, by = "userId") %>%
  mutate(pred = mu + b_i + b_u) %>%
  pull(pred)
model_6_rmse_edx <- RMSE(predicted_ratings_6_edx, edx$rating)
model_6_rmse_edx

```

```
## [1] 0.857032
```

All right, let's get the residuals of the prediction and then perform matrix factorization on the residuals.

```

edx_residual <- edx %>%
  left_join(b_i_reg, by = "movieId") %>%
  left_join(b_u_reg, by = "userId") %>%
  mutate(residual = rating - mu - b_i - b_u) %>%
  select(userId, movieId, residual)
head(edx_residual)

```

```

##   userId movieId residual
## 1      1      122 0.805170
## 2      1      185 0.535750
## 3      1      292 0.247181
## 4      1      316 0.315499
## 5      1      329 0.327707
## 6      1      355 1.176398

```

3. Use the recosystem library to perform the matrix factorization

Now let's use the *recosystem* library to perform the matrix factorization on the residuals. Both training and validation sets need to be organized to 3 columns: user, item (movies), value (ratings or residuals). Then they need to be transformed into matrix format. Next we write these datasets into hard disk, which will later be assigned to *train_set* and *valid_set* to build the "recosystem". A recommender object *r* will be built using *Reco()* in the *recosystem* package and parameters trained using the *train_set*.

Next the parameters will be used to build the prediction model. Because here we are modeling the residuals after Model 6, we add up the base prediction of Model 6 and the residuals predicted here to get the final prediction for the *validation* set. RMSE can be evaluated and compared with previous models.

```

# as matrix
edx_for_mf <- as.matrix(edx_residual)
validation_for_mf <- validation %>%

```

```

    select(userId, movieId, rating)
validation_for_mf <- as.matrix(validation_for_mf)

# write edx_for_mf and validation_for_mf tables on disk
write.table(edx_for_mf , file = "trainset.txt" , sep = " " , row.names = FALSE, col.names = FALSE)
write.table(validation_for_mf, file = "validset.txt" , sep = " " , row.names = FALSE, col.names = FALSE)

# use data_file() to specify a data set from a file in the hard disk.
set.seed(2019)
train_set <- data_file("trainset.txt")
valid_set <- data_file("validset.txt")

# build a recommender object
r <-Reco()

# tuning training set
opts <- r$tune(train_set, opts = list(dim = c(10, 20, 30), lrate = c(0.1, 0.2),
                                     costp_l1 = 0, costq_l1 = 0,
                                     nthread = 1, niter = 10))

opts

```

```

## $min
## $min$dim
## [1] 30
##
## $min$costp_l1
## [1] 0
##
## $min$costp_l2
## [1] 0.01
##
## $min$costq_l1
## [1] 0
##
## $min$costq_l2
## [1] 0.1
##
## $min$lrate
## [1] 0.1
##
## $min$loss_fun
## [1] 0.793797
##
##
## $res
##      dim costp_l1 costp_l2 costq_l1 costq_l2 lrate loss_fun
## 1    10         0     0.01         0     0.01   0.1 0.807220
## 2    20         0     0.01         0     0.01   0.1 0.811662
## 3    30         0     0.01         0     0.01   0.1 0.820999
## 4    10         0     0.10         0     0.01   0.1 0.804869
## 5    20         0     0.10         0     0.01   0.1 0.802334
## 6    30         0     0.10         0     0.01   0.1 0.803364
## 7    10         0     0.01         0     0.10   0.1 0.803780

```

```
## 8 20 0 0.01 0 0.10 0.1 0.795706
## 9 30 0 0.01 0 0.10 0.1 0.793797
## 10 10 0 0.10 0 0.10 0.1 0.824731
## 11 20 0 0.10 0 0.10 0.1 0.823860
## 12 30 0 0.10 0 0.10 0.1 0.823186
## 13 10 0 0.01 0 0.01 0.2 0.809948
## 14 20 0 0.01 0 0.01 0.2 0.822562
## 15 30 0 0.01 0 0.01 0.2 0.837361
## 16 10 0 0.10 0 0.01 0.2 0.805661
## 17 20 0 0.10 0 0.01 0.2 0.804932
## 18 30 0 0.10 0 0.01 0.2 0.807245
## 19 10 0 0.01 0 0.10 0.2 0.802191
## 20 20 0 0.01 0 0.10 0.2 0.799993
## 21 30 0 0.01 0 0.10 0.2 0.799319
## 22 10 0 0.10 0 0.10 0.2 0.823480
## 23 20 0 0.10 0 0.10 0.2 0.822147
## 24 30 0 0.10 0 0.10 0.2 0.820740
```

```
# training the recommender model
r$train(train_set, opts = c(opts$min, nthread = 1, niter = 20))

# Making prediction on validation set and calculating RMSE:
pred_file <- tempfile()
r$predict(valid_set, out_file(pred_file))
predicted_residuals_mf <- scan(pred_file)
predicted_ratings_mf <- predicted_ratings_6 + predicted_residuals_mf
rmse_mf <- RMSE(predicted_ratings_mf, validation$rating) # 0.786256
rmse_results <- bind_rows(rmse_results,
                          data_frame(Model="Matrix Factorization",
                                      RMSE = rmse_mf))
rmse_results
```

```
## # A tibble: 9 x 2
##   Model RMSE
##   <chr> <dbl>
## 1 Just the average 1.06120
## 2 Age Effect Model 1.05239
## 3 Movie Effect Model 0.943909
## 4 Movie + User Effects Model 0.865349
## 5 Regularized Movie Effect Model 0.943852
## 6 Regularized Movie + User Effect Model 0.864818
## 7 Regularized Movie + User Effect Model Version 2 0.864850
## 8 Regularized Movie + User Effect Model Version 3 0.864819
## 9 Matrix Factorization 0.786256
```

Conclusion

From the summarized RMSEs of different models, we can see that matrix factorization largely improved the accuracy of the prediction.

Model	RMSE
Just the average	1.061202
Age Effect Model	1.052393
Movie Effect Model	0.943909
Movie + User Effects Model	0.865349
Regularized Movie Effect Model	0.943852
Regularized Movie + User Effect Model	0.864818
Regularized Movie + User Effect Model Version 2	0.864850
Regularized Movie + User Effect Model Version 3	0.864819
Matrix Factorization	0.786256

MovieLens is a classical dataset for recommendation system and represents a challenge for development of better machine learning algorithm. In this project, the “Just the average” model only gives a RMSE of 1.0612, and the best baseline model (Model 6: Regularized Movie + User Effect Model) could largely improved it to 0.8648. Furthermore, matrix factorization greatly brought it down to 0.7863. In conclusion, matrix factorization appears to be a very powerful technique for recommendation system, which usually contains large and sparse dataset making it hard to make prediction using other machine learning strategies. The effects of age and genres could be further explored to improve the performance of the model. The Ensemble method should also be considered in the future to apply on the MovieLens dataset, in order to combine the advantages of various models and enhance the overall performance of prediction.